
REAL-TIME CASE STUDY #3

E-Commerce Marketplace Orders, Customers & Seller Performance Analytics

BUSINESS CONTEXT

You are a **Data Engineer** working for a **large e-commerce marketplace** (Amazon / Flipkart / Meesho style).

Business wants to:

- Identify **high-value customers**
- Measure **seller performance**
- Understand **category-wise revenue**
- Detect **inactive sellers and customers**
- Optimize **order fulfillment**

Data is coming from **multiple internal systems**, each with **data quality issues**.

Your task is to **clean, integrate, analyze, and optimize** the data using **PySpark**.

DATA SOURCES PROVIDED

You are given **5 raw datasets**.

DATASET 1 – CUSTOMER MASTER (CORRUPTED)

```
raw_customers = [
    ("C001", "Rahul", "29", "Bangalore", "Electronics,Fashion"),
    ("C002", "Sneha", "Thirty Two", "Delhi", "Fashion"),
    ("C003", "Aman", None, "Mumbai", ["Home", "Electronics"]),
    ("C004", "Pallavi", "27", "Pune", "Electronics|Beauty"),
```

```
("C005", "", "35", "Chennai", None)  
]
```

Known Issues

- Age mixed formats
 - Interests as string / array / multiple delimiters
 - Missing names
 - Null interests
-

DATASET 2 – SELLER MASTER

```
raw_sellers = [  
    ("S001", "TechWorld", "Electronics", "2019-06-01"),  
    ("S002", "FashionHub", "Fashion", "01/07/2020"),  
    ("S003", "HomeEssentials", "Home", "2018/09/15"),  
    ("S004", "BeautyStore", "Beauty", "invalid_date")  
]
```

Known Issues

- Seller onboarding date in multiple formats
 - Invalid dates
-

DATASET 3 – PRODUCT CATALOG

```
raw_products = [  
    ("P001", "Laptop", "Electronics", "S001", "55000"),  
    ("P002", "Headphones", "Electronics", "S001", "2500"),  
    ("P003", "T-Shirt", "Fashion", "S002", "1200"),  
    ("P004", "Sofa", "Home", "S003", "45000"),  
    ("P005", "Face Cream", "Beauty", "S004", "800")  
]
```

Known Issues

- Price as string
 - Category repetition
 - Foreign key to seller
-

DATASET 4 – ORDERS DATA

```
raw_orders = [
    ("O001", "C001", "P001", "2024-01-05", "Delivered", "55000"),
    ("O002", "C002", "P003", "05/01/2024", "Cancelled", "0"),
    ("O003", "C003", "P004", "2024/01/06", "Delivered", "45000"),
    ("O004", "C004", "P005", "invalid_date", "Delivered", "800"),
    ("O005", "C001", "P002", "2024-01-10", "Delivered", "2500"),
    ("O006", "C005", "P003", "2024-01-12", "Delivered", "1200")
]
```

Known Issues

- Multiple date formats
 - Invalid dates
 - Cancelled orders
 - Revenue as string
-

DATASET 5 – CUSTOMER ACTIVITY LOGS

```
raw_activity = [
    ("C001", "search,view,add_to_cart", {"device": "mobile"}, 180),
    ("C002", ["search", "view"], "device=laptop", 90),
    ("C003", "search|view|purchase", None, 120),
    ("C004", None, {"device": "tablet"}, 60),
    ("C005", "search", {"device": "mobile"}, 30)
]
```

Known Issues

- Actions in multiple formats
- Metadata as JSON-like strings
- Missing actions

BUSINESS QUESTIONS TO ANSWER

You must complete **all sections**.

PART A – DATA CLEANING & STRUCTURING

1. Design **explicit schemas** for all datasets
 2. Normalize:
 - Age
 - Prices
 - Dates
 3. Convert interests and actions into arrays
 4. Handle missing and invalid records gracefully
 5. Produce clean DataFrames:
 - `customers_df`
 - `sellers_df`
 - `products_df`
 - `orders_df`
 - `activity_df`
-

PART B – DATA INTEGRATION (JOINS)

6. Join orders with products
 7. Join products with sellers
 8. Join orders with customers
 9. Decide which table(s) should be **broadcast**
 10. Prove your decision using `explain(True)`
 11. Eliminate orphan records
-

PART C – ANALYTICS & AGGREGATIONS

-
- 12. Total revenue per category
 - 13. Total revenue per seller
 - 14. Total orders per customer
 - 15. Average order value per customer
 - 16. Identify sellers with **zero delivered orders**
-

PART D – WINDOW FUNCTIONS

- 17. Rank customers by total spend (overall)
 - 18. Rank sellers by revenue **within each category**
 - 19. Calculate running revenue per day
 - 20. Identify top 2 products per category by revenue
-

PART E – UDF (ONLY IF REQUIRED)

- 21. Classify customers into spending tiers:

- o High
- o Medium
- o Low

Rules:

- Prefer built-in functions
 - Use UDF only if unavoidable
 - Justify your choice
-

PART F – SORTING & ORDERING

- 22. Sort categories by total revenue (descending)
 - 23. Sort sellers by revenue within category
 - 24. Explain why sorting caused a shuffle
-

PART G – SET OPERATIONS

Create two DataFrames:

- Customers who **placed orders**
- Customers who **were active (search/view)**

25. Find customers who were active but never ordered
 26. Find customers who ordered and were active
 27. Explain why set operations differ from joins
-

PART H – DAG & PERFORMANCE ANALYSIS

28. Run `explain(True)` for:

- Product → Seller join
- Window ranking
- Sorting

29. Identify:

- Shuffles
- Broadcast joins
- Sort stages

30. Suggest one performance improvement
