

PYSPARK EXERCISES – LARGE DATASET

Dataset: `orders_large_bad.csv` /
`orders_large_bad.json`

CONTEXT

You are a **Data Engineer** working with **raw production data** ingested from multiple upstream systems.

The dataset:

- Is large (hundreds of thousands of rows)
- Has bad formatting
- Has data quality issues
- Is NOT analytics-ready

Your job is to **ingest, clean, optimize, analyze, and store** this data using PySpark.

INPUT DATA

Files provided:

- `orders_large_bad.csv`
- `orders_large_bad.json`

Columns:

order_id
customer_id
city
category
product
amount
order_date
status

◆ PHASE 1 – INGESTION & FIRST INSPECTION

Exercises

1. Read the CSV file into a DataFrame
 2. Disable schema inference and read everything as string
 3. Print schema and record count
 4. Display 20 random rows
 5. Identify **at least 5 data quality issues** by observation
 6. Read the JSON file and compare schema and row count with CSV
-

◆ PHASE 2 – SCHEMA ENFORCEMENT & VALIDATION

Exercises

7. Define an explicit schema using `StructType`
 8. Re-read the CSV using the defined schema
 9. Identify rows that fail schema expectations
 10. Explain why schema inference is dangerous at scale
-

◆ PHASE 3 – STRING CLEANING & STANDARDIZATION

Exercises

11. Trim leading and trailing spaces from all string columns
12. Standardize `city`, `category`, and `product` values

-
- 13. Convert all categorical columns to a consistent case
 - 14. Identify how many distinct city values existed before vs after cleaning
-

◆ PHASE 4 – AMOUNT CLEANING (CRITICAL)

Exercises

- 15. Identify invalid values in the `amount` column
 - 16. Remove commas from numeric strings
 - 17. Convert `amount` to IntegerType safely
 - 18. Handle empty, null, and invalid values explicitly
 - 19. Count how many records were affected during amount cleaning
-

◆ PHASE 5 – DATE PARSING & NORMALIZATION

Exercises

- 20. Identify all date formats present in `order_date`
 - 21. Parse valid dates into `DateType`
 - 22. Handle invalid dates gracefully
 - 23. Create a clean `order_date_clean` column
 - 24. Count records with invalid dates
-

◆ PHASE 6 – BUSINESS FILTERING & DEDUPLICATION

Exercises

- 25. Identify duplicate `order_id` values
- 26. Remove duplicate orders safely

-
- 27. Keep only records with status = `Completed`
 - 28. Validate record counts before and after filtering
-

◆ PHASE 7 – PERFORMANCE & PARTITION AWARENESS

Exercises

- 29. Check the default number of partitions
 - 30. Run a heavy `groupBy` and observe execution time
 - 31. Use `explain(True)` to identify shuffle stages
 - 32. Repartition the DataFrame by `city`
 - 33. Compare execution plans before and after repartition
-

◆ PHASE 8 – ANALYTICS ON LARGE DATA

Exercises

- 34. Calculate total revenue per city
 - 35. Calculate total revenue per category
 - 36. Calculate total revenue per product
 - 37. Identify top 10 products by revenue
 - 38. Calculate average order value per city
-

◆ PHASE 9 – WINDOW FUNCTIONS (BIG DATA SAFE)

Exercises

- 39. Rank cities by total revenue
- 40. Rank products within each category by revenue

-
- 41. Identify the top product per category
 - 42. Identify top 3 cities using window functions
-

◆ PHASE 10 – CACHING & REUSE

Exercises

- 43. Identify DataFrames reused multiple times
 - 44. Apply caching strategically
 - 45. Re-run analytics and observe performance
 - 46. Unpersist when cache is no longer needed
 - 47. Explain why over-caching is dangerous
-

◆ PHASE 11 – FILE FORMAT STRATEGY

Exercises

- 48. Write the cleaned order-level dataset to **Parquet**
 - 49. Partition the Parquet output by `city`
 - 50. Write aggregated analytics to **ORC**
 - 51. Read both formats back and validate schema
 - 52. Compare number of output files generated
-

◆ PHASE 12 – DEBUGGING & FAILURE

SCENARIOS

Exercises

- 53. Explain why the following line breaks pipelines:

```
df = df.filter(df.amount > 50000).show()
```

-
- 54. Create a scenario that produces a `NoneType` error
 - 55. Identify a transformation that causes a wide shuffle
 - 56. Explain how you would debug a slow Spark job
-

◆ PHASE 13 – FINAL VALIDATION

Exercises

- 57. Validate no nulls in critical columns
 - 58. Confirm correct data types for all columns
 - 59. Validate final record count
 - 60. Document three optimization decisions you made
-